

A.23 Monetary-policy sensitivity: data and impulse responses

Shock series. We use the high-frequency identified monetary-policy surprises of [Jarocinski and Karadi \(2020\)](#), who decompose FOMC announcement-window co-movements of federal funds futures and the S&P 500 into a “pure” monetary-policy component and a central-bank information component. We use the monthly pure-policy median series (MP_median) and aggregate to the quarterly frequency as the in-quarter sum. The resulting series is positive on net tightenings and negative on net easings and is, by construction, orthogonal to information shocks. For the main-text regression we use shocks at $t \in 2015\text{Q1}–2021\text{Q3}$ so that the forward one-quarter outcome window closes no later than 2021Q4; for the local projections we restrict each horizon h to $t + h \leq 2021\text{Q4}$.

Sample. The LinkedIn/Revelio Labor firm-quarter panel covers 284,384 firms observed 2015Q1–2024Q4. Firm-level adoption (“Adopter_{*i*}”) is defined as in Section 4.2.2. The main-text outcomes are the forward one-quarter log-change in total headcount ($\log L_{i,t+1} - \log L_{i,t}$) and in junior headcount (seniority levels 1–2; $\log L_{i,t+1}^{1,2} - \log L_{i,t}^{1,2}$), enforced by a strict one-calendar-quarter gap and winsorised per outcome at the 0.5/99.5% level. The local-projection outcomes below use the cumulative log-change $\log L_{i,t+h} - \log L_{i,t-1}$ at each horizon h .

Local-projection impulse responses. To verify that the null result in Table 3 is not horizon-specific, we estimate the lag-augmented Jordà local projection ([Jordà, 2005](#); [Montiel Olea and Plagborg-Møller, 2021](#))

$$\log L_{i,t+h} - \log L_{i,t-1} = \alpha_i^h + \gamma_t^h + \beta_h(\text{Adopter}_i \times \varepsilon_t^m) + \sum_{\ell=1}^{12} \theta_{h,\ell} \Delta \log L_{i,t-\ell} + \eta_{i,t}^h, \quad (35)$$

for horizons $h = 0, \dots, 8$ quarters. The specification includes 12 quarterly lags of log-employment growth as controls, following [Montiel Olea and Plagborg-Møller \(2021\)](#), and firm and year-quarter fixed effects; standard errors are Cameron–Gelbach–Miller two-way cluster-robust at the firm and year-quarter level. We also report the pooled impulse response from the analogous specification without the adopter interaction and with firm fixed effects only (time fixed effects would absorb ε_t^m).

Figure A.25 plots the pooled (top row) and adopter-interaction (bottom row) impulse responses with 95% confidence bands. The interaction β_h is statistically indistinguishable from zero at every horizon for both total and junior employment, consistent with the one-quarter sensitivity estimate in the main text.

Lag-augmented LP-IRFs (MOPM 2021; L=12 qtrly lags of $\Delta \log L$): LinkedIn employment · JK(2020) MP_pure shock · 2015Q1–2021Q4 (pre-2022)




Figure A.25: Impulse responses of LinkedIn employment to Jarocinski and Karadi (2020) pure monetary-policy shocks, lag-augmented Jordà local projection, pre-2022 sample. Top row: pooled response β_h from $\log L_{i,t+h} - \log L_{i,t-1} = \alpha_i^h + \beta_h \epsilon_t^m + \text{lag controls} + \eta$, firm fixed effects only. Bottom row: adopter-interaction response β_h from equation (35), firm and year-quarter fixed effects. Shaded bands are 95% confidence intervals using Cameron–Gelbach–Miller two-way cluster-robust standard errors at the firm and year-quarter level.

A.24 Remote Work as a Confounder

This appendix provides details on the remote-work robustness exercise discussed in Section 5.4. The concern, raised by Lambert and Schindler (2026), is that the post-ChatGPT decline in junior hiring at GenAI adopters may instead reflect the post-pandemic normalization of remote work. Remote-friendly occupations expanded sharply during the pandemic and contracted as firms scaled back remote arrangements. In addition, work-from-home arrangements may reduce demand for junior workers by increasing supervision costs and weakening opportunities for on-the-job learning. Because occupations that are amenable to remote work overlap substantially with those exposed to GenAI, a reallocation away from remote-able jobs could, in principle, masquerade as a GenAI effect.

Data. We measure the remote suitability of each position using the `remote_suitability` score in the Revelio Labs individual-positions data, a continuous index of how amenable a position’s occupation is to remote work. We download this score for the positions of our sample firms and merge it onto the same position panel used to construct our employment measures, so that the resulting counts line up one-for-one with the seniority head counts used throughout the paper. Mirroring our treatment of the Eloundou et al. (2024) GenAI-exposure measure, we classify individual positions into high- and low-remote-suitability groups using the top and bottom quartiles of the position-level score distribution; the middle two quartiles are excluded from the split. For each firm-month we then count the number of junior and senior workers in high- and low-remote-suitability positions, and aggregate to the firm-quarter level.

Specification. Within each remote-suitability group separately, we estimate the seniority triple-difference of Equation 11, comparing log employment of junior versus senior workers across adopting and non-adopting firms over time. Concretely, we run the specification once restricted to high-remote-suitability positions (junior versus senior head counts in high-remote jobs) and once restricted to low-remote-suitability positions, in each case absorbing firm-by-quarter, industry-by-quarter-by-seniority, and industry-by-seniority fixed effects, with standard errors clustered at the firm level. Coefficients are normalized to zero in 2022Q4, the quarter before ChatGPT’s wide diffusion.

Result. Figure A.26 plots the triple-difference coefficients for the two groups. The key comparison is the low-remote-suitability series: these are jobs that cannot be performed from home, so a remote-work correction should leave them unaffected. Instead, the seniority-biased junior decline at adopters appears after ChatGPT in *both* groups and is of broadly similar magnitude. Because a remote-work correction operates

primarily on the remote-able margin, it cannot explain a junior decline that is also present—at comparable magnitude—among low-remote-suitability positions. This indicates that our main results are not driven by the post-pandemic reallocation of remote work.




Figure A.26: Seniority Triple-Difference, Separately Within High- and Low-Remote-Suitability Positions

Notes: Each series plots the junior-relative-to-senior triple-difference coefficients (Equation 11), estimated separately using employment in high- and low-remote-suitability positions. Remote suitability is measured by the Revelio Labs `remote_suitability` score; high (low) jobs are those in the top (bottom) quartile of the score distribution. The dark solid line corresponds to high-remote-suitability positions and the light dashed line to low-remote-suitability positions; shaded bands are 95% confidence intervals based on standard errors clustered at the firm level. Coefficients are normalized to zero in 2022Q4. The vertical line marks the launch of ChatGPT (2022Q4).