

Generative AI and Labor Market Outcomes: Evidence from the United Kingdom

Bouke Klein Teeselink^{*1,2,3}

¹*Department of Political Economy, King's College London*

²*AI Objectives Institute*

³*King's Institute for Artificial Intelligence*

December 21, 2025

Abstract

This paper examines the effects of LLMs on UK labor market outcomes. We use a difference-in-differences design that compares outcomes across firms and occupations from 2021 to 2025 based on their differential exposure to LLM capabilities. Our results show that highly exposed firms reduce employment, with effects concentrated in junior positions. These firms sharply curtail new hiring, with technical and creative roles experiencing the steepest declines. At the occupation level, roles more exposed to LLMs also show substantial reductions in job listings. The observed effects concentrate almost entirely in high-wage segments. Hence, our findings suggest that LLMs may compress wage inequality.

Keywords: Generative AI; Large Language Models; Automation; Labor demand; United Kingdom.

JEL: O33; J23; J31

*Contact: bouke.klein.teeselink@kcl.ac.uk. I would like to thank Philip Tomei, Barbara Piotrowska, Bharat Chandar, and seminar participants of the Economics of TAI group for insightful comments. I particularly want to thank Jacob Schaal for exceptional research assistance.

Table 3: Effect of LLM exposure on firm-level job outcomes

	Employees	Low seniority	High seniority	Frac. low seniority	Av. comp.
LLM exposure x Post	−0.003* (0.002)	−0.004** (0.002)	−0.002 (0.002)	−0.001** (0.0003)	106.313* (59.029)
Firm fixed effects	Yes	Yes	Yes	Yes	Yes
Industry-by-period FE	Yes	Yes	Yes	Yes	Yes
Observations	1,268,013	1,267,955	1,266,331	1,267,955	1,268,013
Adj. R-squared	0.983	0.983	0.981	0.979	0.978

Notes: This table reports difference-in-differences estimates of the effect of LLM exposure on firm-level employment outcomes, averaged over the post-treatment period. *Employees* is the total number of jobs in the firm. *Low seniority* and *High seniority* denote the number of positions in seniority levels 1–3 and 4–7, respectively. *Low seniority (%)* is the share of jobs that are low-seniority. *Av. comp.* is the mean annual salary plus bonuses. All regressions include firm fixed effects and industry-by-period fixed effects. Standard errors are clustered at the firm level and reported in parentheses. Significance is denoted by * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The regression results in Table 3 show, averaged over the entire post-treatment period, a significant 0.4 percent reduction in junior employees per standard deviation increase in exposure, and a statistically insignificant reduction in senior employees. The fraction of junior employees goes down by 0.1 percentage points per standard deviation of LLM exposure ($p = 0.028$).

The observed compositional shift towards more senior roles has direct implications for compensation patterns. Figure 3 shows that average compensation rises modestly in highly exposed firms following the introduction of ChatGPT, with a one standard deviation increase in exposure associated with a £106 increase in average compensation, although the overall estimate is only marginally significant.

4.2 Job Posting Dynamics

4.2.1 Firm-Level Postings

While employment stocks adjust slowly, job postings provide a real-time window into firms’ evolving labor demand. We therefore examine how firms’ hiring intentions respond to LLM exposure by analyzing whether exposed firms post fewer vacancies overall and whether they adjust the composition of roles they seek to fill. This posting data captures the extensive margin of recruitment activity and reveals forward-looking decisions that precede the employment stock changes documented above. By tracking the same firms across both datasets, we can directly observe

Table 4: Effect of LLM exposure on occupation-level job postings

	Job postings (log)	Salary	Salary (log)	Salary (SD)
LLM exposure x Post	-0.039** (0.018)	-495.483*** (135.270)	-0.011*** (0.003)	-56.802 (108.358)
Occupation fixed-effects	Yes	Yes	Yes	Yes
Year-by-month fixed effects	Yes	Yes	Yes	Yes
Observations	16,815	16,815	16,815	16,804
Adjusted R ²	0.962	0.948	0.953	0.729

Notes: This table reports difference-in-differences estimates of the effect of LLM exposure on occupation-level outcomes. *Job postings (log)* is the log number of new vacancies in occupation o in month t . *Salary* is the mean advertised annual compensation, *Salary (log)* is its logarithm, and *Salary (SD)* is the within-occupation standard deviation of posted salaries. *Exposure* is the occupation-level LLM exposure score from Eloundou et al. (2024). *Post* equals one for November 2022 onward. Standard errors are clustered at the occupation level. Significance is denoted by * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

4.2.2 Occupation-Level Postings

This analysis examines occupation-level job posting data to understand how aggregate labor demand evolves for specific roles. This analysis exploits variation in LLM exposure across 384 occupations, from minimally affected roles such as cooks (exposure score: 0) to highly exposed positions like proofreaders and copy markers (exposure score: 0.86). We examine three outcomes: posting volumes to measure aggregate demand, advertised salaries to assess compensation responses, and within-occupation salary dispersion to measure wage polarization.

Figure 6 presents event study estimates for job posting volumes. Prior to November 2022, high- and low-exposure occupations follow parallel trajectories, validating our identification strategy. Following ChatGPT’s release, however, exposed occupations experience a sustained reduction in hiring demand, with effects first becoming statistically significant after 7 months and intensifying thereafter. After 20 months, we see a slight resurgence, and job postings converge back to their original levels. When averaged across the post-treatment period (Table 4), a one standard deviation increase in LLM exposure is associated with a 3.9% reduction in posting volume ($p = 0.029$). This is an economically meaningful change.

Turning to compensation, we observe immediate and persistent wage pressure for exposed occupations. Figure 7 indicates that posted salaries begin declining shortly after ChatGPT’s introduction. A one standard deviation increase in LLM exposure is associated with a £495 decline in posted salaries and a 1.1% reduction in log

ChatGPT’s release, with a one standard deviation increase in LLM exposure reducing the probability of posting any vacancy by 1.0 percentage point. This reduction emerges gradually and becomes significant after approximately 10 months, mirroring the timeline of employment stock adjustments. When we examine the composition of roles driving this reduction, we find that technical positions and creative roles experience the sharpest declines. Sales representative positions, by contrast, show a statistically significant increase, suggesting that customer-facing roles requiring direct interpersonal interaction remain robust to LLM-based automation.

In addition to within-firm dynamics, our occupation-level analysis reveals how labor demand shifts for entire job categories regardless of firm exposure. A one standard deviation increase in LLM exposure is associated with a 3.9 percent reduction in posting volumes, accompanied by £495 (1.1 percent) salary declines. The apparent inconsistency of declining posted salaries for exposed occupations alongside increasing firm-level compensation shows how aggregation could mask opposing forces. While LLM exposure drives down market wages for specific occupations, within-firm rebalancing toward more senior positions pushes average firm-level wages in the opposite direction.

Our heterogeneity analysis reveals that the observed labor market disruptions are almost entirely driven by high-remuneration segments of the labor market. This concentration of effects in higher-wage segments contradicts historical automation patterns that predominantly displaced routine middle-skill work, and suggests that LLMs might lead to a compression of the wage distribution.

Our findings carry important implications for understanding technological displacement and designing policy responses. The concentration of job losses among entry-level positions disrupts traditional skill development pathways where workers master progressively complex tasks through hands-on experience. Without junior roles serving as training grounds, firms may struggle to develop senior talent internally while new entrants face shrinking opportunities to gain experience (Ide, 2025). The freeze in hiring followed by employment stock changes suggests that forward-looking indicators like job postings may provide early warning signals for policymakers. The sharp reductions in technical and analytical roles alongside resilience in sales positions indicates that human interaction and relationship-building skills remain valuable complements to LLM capabilities, suggesting that retraining programs should emphasize these interpersonal competencies. Policy responses