

Generative AI and Labor Market Outcomes: Evidence from the United Kingdom

Bouke Klein Teeselink^{*1,2,3}

¹*Department of Political Economy, King's College London*

²*AI Objectives Institute*

³*King's Institute for Artificial Intelligence*

December 21, 2025

Abstract

This paper examines the effects of LLMs on UK labor market outcomes. We use a difference-in-differences design that compares outcomes across firms and occupations from 2021 to 2025 based on their differential exposure to LLM capabilities. Our results show that highly exposed firms reduce employment, with effects concentrated in junior positions. These firms sharply curtail new hiring, with technical and creative roles experiencing the steepest declines. At the occupation level, roles more exposed to LLMs also show substantial reductions in job listings. The observed effects concentrate almost entirely in high-wage segments. Hence, our findings suggest that LLMs may compress wage inequality.

Keywords: Generative AI; Large Language Models; Automation; Labor demand; United Kingdom.

JEL: O33; J23; J31

*Contact: bouke.klein.teeselink@kcl.ac.uk. I would like to thank Philip Tomei, Barbara Piotrowska, Bharat Chandar, and seminar participants of the Economics of TAI group for insightful comments. I particularly want to thank Jacob Schaal for exceptional research assistance.

in exposed occupations. While their analysis covers 25 broad occupation groups, we exploit variation across a much wider variety of job types, as well as a firm-level analysis that examines within-firm dynamics that would be invisible in occupational aggregates. In addition, our individual-level posting data enables examination of a wider set of outcome variables such as salary distributions and seniority-specific employment, as well as heterogeneity in treatment effects by pre-treatment compensation or location.

Taken together, our findings raise important policy questions. The disproportionate elimination of entry-level positions disrupts the conventional model of skill development, where employees master progressively complex tasks through hands-on experience with routine work. Without junior roles that serve as training grounds, firms may struggle to develop future senior talent internally, creating a hollowed-out labor market where entry becomes increasingly difficult (Ide, 2025). This dynamic could exacerbate inequality as those with pre-existing experience benefit from reduced competition while new entrants face shrinking opportunities. Policy responses must address both immediate displacement and longer-term skill development pathways.

2 Data

Our primary data comes from Revelio Labs, which aggregates and processes publicly available employment information from professional networking platforms, company websites, and job boards using machine learning algorithms to create proprietary longitudinal labor market databases. Revelio Labs transforms semi-structured public profile information into a structured panel dataset through systematic parsing, standardization, and cleaning procedures. Their data have been validated against official statistics and used in recent academic research on labor market dynamics (Fadhel et al., 2021; Eisfeldt et al., 2023; Hampole et al., 2025). In Appendix A4, we separately validate that Revelio Labs employment and job posting data closely track Office for National Statistics aggregates both over time and across industries. We use two separate Revelio Labs datasets: work profiles and job postings.²

²The data used in this paper were downloaded on 12 August 2025.

2.1 Job Profile Data

The first dataset consists of all publicly available professional LinkedIn profiles between January 2021 through June 2025. The dataset includes approximately 75 million employment spells from 23 million unique individuals across 2.5 million UK-based firms. Each observation includes the worker’s job title, employer, start and end dates (if available), seniority level,³ and compensation.⁴ The data provide near-real-time employment stocks, which enable us to track within-firm workforce composition changes over time.

From the panel of individual LinkedIn profiles, we construct a monthly firm-level panel from January 2021 through June 2025. For each firm-month, we calculate total employment as the count of unique individuals with active employment spells at a certain company in that month, and log-transform this variable. This is our main employment stock measure. We then separate total employment into junior and senior employment stocks based on Revelio seniority classifications of jobs. We also calculate average compensation as the log of the mean of reported salaries plus bonuses in a firm. We only focus on firms that remained in business throughout the sample period.

To calculate firm-level exposure scores, we use the assessment framework developed by Eloundou et al. (2024). They evaluate 19,265 detailed work activities and tasks from the O*NET 27.2 database to determine whether LLM assistance could reduce task completion time by at least 50% while maintaining quality. We focus on their ‘E1’ measure, which captures exposure to current LLMs, rather than LLMs combined with additional software. We do so because it is the most appropriate metric for identifying short-term labor market effects that do not require additional software development or complex organizational changes. To obtain our firm-level

³Revelio Labs determines seniority using an ensemble model that combines three components: (i) current job characteristics including title, company, and industry; (ii) employment history including tenure and previous position seniority; and (iii) individual age. These three scores are averaged to create a continuous seniority metric, which is then mapped to seven ordinal levels based on keyword benchmarks (e.g., “junior,” “senior,” “director”). The levels range from Entry Level (e.g., interns, trainees) through Junior, Associate, Manager, and Director levels, to Executive and Senior Executive levels (e.g., C-suite positions).

⁴Revelio Labs estimates compensation using a predictive model trained on publicly available visa application data, self-reported salaries, and job posting information. The model incorporates position-specific variables (job title, seniority level, company, location), individual-specific factors (tenure at company), and temporal controls (year of observation) to generate salary predictions for positions where compensation data are not directly observed.

Table 4: Effect of LLM exposure on occupation-level job postings

	Job postings (log)	Salary	Salary (log)	Salary (SD)
LLM exposure x Post	-0.039** (0.018)	-495.483*** (135.270)	-0.011*** (0.003)	-56.802 (108.358)
Occupation fixed-effects	Yes	Yes	Yes	Yes
Year-by-month fixed effects	Yes	Yes	Yes	Yes
Observations	16,815	16,815	16,815	16,804
Adjusted R ²	0.962	0.948	0.953	0.729

Notes: This table reports difference-in-differences estimates of the effect of LLM exposure on occupation-level outcomes. *Job postings (log)* is the log number of new vacancies in occupation o in month t . *Salary* is the mean advertised annual compensation, *Salary (log)* is its logarithm, and *Salary (SD)* is the within-occupation standard deviation of posted salaries. *Exposure* is the occupation-level LLM exposure score from Eloundou et al. (2024). *Post* equals one for November 2022 onward. Standard errors are clustered at the occupation level. Significance is denoted by * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

4.2.2 Occupation-Level Postings

This analysis examines occupation-level job posting data to understand how aggregate labor demand evolves for specific roles. This analysis exploits variation in LLM exposure across 384 occupations, from minimally affected roles such as cooks (exposure score: 0) to highly exposed positions like proofreaders and copy markers (exposure score: 0.86). We examine three outcomes: posting volumes to measure aggregate demand, advertised salaries to assess compensation responses, and within-occupation salary dispersion to measure wage polarization.

Figure 6 presents event study estimates for job posting volumes. Prior to November 2022, high- and low-exposure occupations follow parallel trajectories, validating our identification strategy. Following ChatGPT’s release, however, exposed occupations experience a sustained reduction in hiring demand, with effects first becoming statistically significant after 7 months and intensifying thereafter. After 20 months, we see a slight resurgence, and job postings converge back to their original levels. When averaged across the post-treatment period (Table 4), a one standard deviation increase in LLM exposure is associated with a 3.9% reduction in posting volume ($p = 0.029$). This is an economically meaningful change.

Turning to compensation, we observe immediate and persistent wage pressure for exposed occupations. Figure 7 indicates that posted salaries begin declining shortly after ChatGPT’s introduction. A one standard deviation increase in LLM exposure is associated with a £495 decline in posted salaries and a 1.1% reduction in log