


► 5 Adjustment of survey sub-sample

In the next step, we compared the average of expert scores with the original mean survey scores using two independent AI models as impartial arbiters to propose an adjusted score.¹³ The models were informed about the sources of the scores in such a way that would prevent them from automatically siding with the expert group: group one was defined as *“practical knowledge of the work and types of tasks they were presented, since they worked in the same occupational groups from which we selected tasks they scored”*, while the second group was described as *“respondents with more theoretical knowledge of technology and labour markets”*. Rather than simply averaging the two scores, the models were explicitly instructed to focus their analysis on the task characteristics and the current capabilities of generative AI and required to provide justifications for each adjustment and its relative position to the original scores.

► Figure 13. Task-level scores from the survey and experts, compared to AI-arbitrated scores




The overall correlation between the adjusted scores produced by the two AI models is 0.96, with correlations within occupational groups exceeding 0.9 in all but one case¹⁴ (Figure 14). This high level of agreement demonstrates that, when provided with sufficient background knowledge in the form of reference ranges of scores, the models converge on remarkably similar predictions. Notably, only three tasks exhibited an absolute difference in scores exceeding 0.2 (orange dots in Figure 13). In these cases, GPT-4o tended to adopt a more optimistic perspective on the feasibility of automation, whereas Gemini emphasized the importance of human judgment as a critical factor.

¹³ We used OpenAI's GPT-4o and Google's Gemini Flash 1.5, which have comparable abilities. We also tested OpenAI o1, which at the time of writing was among the models with strongest capacities of « algorithmic reasoning » and observed very similar results. The process was set up as a cycle of sequential API calls, which loop over each task and capture the response into a data frame.

¹⁴ Public authorities, senior officials and managers in Figure 10.

► Figure 14. Comparison of Adjusted Scores: GPT-4o vs Gemini by Occupational Group



For the final adjustment, we averaged the scores of the two arbitrating models and manually reviewed the justifications, without observing any significant inconsistencies. Table 4 illustrates the result of the adjustment process through examples of tasks where expert scores were significantly higher or lower than the survey averages, with the final adjusted score and GPT-4o justification. For tasks such as “operating injection molding machines” or “performing mechanical processing of silicon”, experts highlighted the limitations of generative AI in handling physical and manual components, leading to significantly reduced adjusted scores compared to the survey. Conversely, for tasks like “maintaining documentation of classes with children” or “performing financial estimates for household budgets”, the experts recognized the structured and analytical nature of these tasks, which generative AI can assist with to a substantial degree.

► Table 4. Sample of adjustments and justifications between the survey and expert scores (examples of largest upward and downward revisions)

Task				Justification
Applying various welding methods and using different types of welding machines and equipment, in accordance with training or qualifications.	0.43	0.05	0.05	The task involves hands-on physical work requiring skill and precision, such as welding using various machines, which is not currently automatable by generative AI technologies like ChatGPT. These technologies excel in tasks involving data processing, language understanding, and information generation but do not possess capabilities for physical task execution, hence aligning closer with the revised score.