

NBER WORKING PAPER SERIES

DIFFERENCE-IN-DIFFERENCES WITH VARIATION IN TREATMENT TIMING

Andrew Goodman-Bacon

Working Paper 25018

<http://www.nber.org/papers/w25018>

NATIONAL BUREAU OF ECONOMIC RESEARCH

1050 Massachusetts Avenue

Cambridge, MA 02138

September 2018

I thank Michael Anderson, Martha Bailey, Marianne Bitler, Brantly Callaway, Kitt Carpenter, Eric Chyn, Bill Collins, John DiNardo, Andrew Dustan, Federico Gutierrez, Brian Kovak, Emily Lawler, Doug Miller, Sayeh Nikpay, Pedro Sant'Anna, Jesse Shapiro, Gary Solon, Isaac Sorkin, Sarah West, and seminar participants at the Southern Economics Association, ASHEcon 2018, the University of California, Davis, University of Kentucky, University of Memphis, University of North Carolina Charlotte, the University of Pennsylvania, and Vanderbilt University. All errors are my own. The views expressed herein are those of the author and do not necessarily reflect the views of the National Bureau of Economic Research.

NBER working papers are circulated for discussion and comment purposes. They have not been peer-reviewed or been subject to the review by the NBER Board of Directors that accompanies official NBER publications.

© 2018 by Andrew Goodman-Bacon. All rights reserved. Short sections of text, not to exceed two paragraphs, may be quoted without explicit permission provided that full credit, including © notice, is given to the source.

In contrast to our substantial understanding of the canonical 2x2 DD model, we know relatively little about the two-way fixed effects DD model when treatment timing varies. We do not know precisely how it compares mean outcomes across groups.⁵ We typically rely on general descriptions of the identifying assumption like “interventions must be as good as random, conditional on time and group fixed effects” (Bertrand, Duflo, and Mullainathan 2004, p. 250), and consequently lack well-defined strategies to test the validity of the DD design with timing. We have limited understanding of the treatment effect parameter that regression DD identifies. Finally, we often cannot evaluate when alternative specifications will work or why they change estimates.⁶

This paper shows that the two-way fixed effects DD estimator in (2) is a weighted average of all possible 2x2 DD estimators that compare timing groups to each other. Some use units treated at a particular time as the treatment group and untreated units as the control group. Some compare units treated at two different times, using the later group as a control before its treatment begins and then the earlier group as a control after its treatment begins. As in any least squares estimator, the weights on the 2x2 DD’s are proportional to group sizes *and* the variance of the treatment dummy within each pair. Treatment variance is highest for groups treated in the middle of the panel and lowest for groups treated at the extremes. This result clarifies the theoretical interpretation and identifying assumptions of the general DD model and creates simple new tools for describing the design and analyzing problems that arise in practice.

By decomposing the DD estimator into its sources of variation (the 2x2 DD’s) and providing an explicit interpretation of the weights in terms of treatment variances, my results

⁵ Imai, Kim, and Wang (2018) note “It is well known that the standard DiD estimator is numerically equivalent to the linear two-way fixed effects regression estimator if there are two time periods and the treatment is administered to some units only in the second time period. Unfortunately, this equivalence result does not generalize to the multi-period DiD design...Nevertheless, researchers often motivate the use of the two-way fixed effects estimator by referring to the DiD design (e.g., Angrist and Pischke, 2009).”

⁶ This often leads to sharp disagreements. See Neumark, Salas, and Wascher (2014) on unit-specific linear trends, Lee and Solon (2011) on weighting and outcome transformations, and Shore-Sheppard (2009) on age time fixed effects.

extend recent research on DD models with heterogeneous effects.⁷ Assuming equal counterfactual trends, Abraham and Sun (2018), Borusyak and Jaravel (2017), and de Chaisemartin and D’Haultfœuille (2018b) show that two-way fixed effects DD yields an average of treatment effects across all groups and times, some of which may have negative weights. My results show how these weights arise from differences in timing and thus treatment variances, facilitating a connection between models of treatment allocation and the interpretation of DD estimates.⁸ I also explain why the negative weights occur: when already-treated units act as controls, *changes* in their treatment effects over time get subtracted from the DD estimate. This negative weighting only arises when treatment effects vary over time, in which case it typically biases regression DD estimates away from the sign of the true treatment effect. This does not imply a failure of the underlying *design*, but it does caution against the use of a single-coefficient two-way fixed effects specification to summarize time-varying effects.

I also show that because regression DD uses group sizes and treatment variances to weight up simple estimates that each rely on common trends between two groups, its identifying assumption is a variance-weighted version of common trends between all groups. The extent to which a group’s differential trend biases the overall estimate equals the difference between how much weight it gets when it acts as the treatment group and how much weight it gets when it acts as the control group. When the earliest and/or latest treated units have low treatment variance, they can get *more* weight as controls than treatments. In designs without untreated units they always

⁷ Early research in this area made specific observations about stylized specifications such as models with no unit fixed effects (Bitler, Gelbach, and Hoynes 2003), or it provided simulation evidence (Meer and West 2013).

⁸ Related results on the weighting of heterogeneous treatment effects does not provide this intuition. Abraham and Sun (2018, p 9) describe the weights in a DD estimate with constant treatment effects as “residual[s] from predicting treatment status, $D_{i,t}$ with unit and time fixed effects.” de Chaisemartin and D’Haultfœuille (2018b, p 7) and Borusyak and Jaravel (2017) describe these same weights as coming from an auxiliary regression and Borusyak and Jaravel (2017, p 10-11) note that “a general characterization of [the weights] does not seem feasible.” Athey and Imbens (2018) also decompose the DD estimator and develop design-based inference methods for this setting. Strezhnev (2018) expresses $\hat{\beta}^{DD}$ as an unweighted average of DD-type terms across pairs of observations and periods.