

Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects*

Liyang Sun[†] and Sarah Abraham[‡]

September 22, 2020

Abstract

To estimate the dynamic effects of an absorbing treatment, researchers often use two-way fixed effects regressions that include leads and lags of the treatment. We show that in settings with variation in treatment timing across units, the coefficient on a given lead or lag can be contaminated by effects from other periods, and apparent pretrends can arise solely from treatment effects heterogeneity. We propose an alternative estimator that is free of contamination, and illustrate the relative shortcomings of two-way fixed effects regressions with leads and lags through an empirical application.

Keywords: difference-in-differences, two-way fixed effects, pretrend test

*We are grateful to Isaiah Andrews, Amy Finkelstein, Anna Mikusheva, and Heidi Williams for their guidance and support. We thank Alberto Abadie, Jonathan Cohen, Nathan Hendren, Peter Hull, Guido Imbens, Yunan Ji, Sylvia Klosin, Kevin Kainan Li, Paichen Li, Therese A. McCarty, Whitney Newey, James Poterba, Pedro H. C. Sant’Anna, Gergely Ujhe-lyi and Helen Willis for helpful discussions. This research was supported by the National Institute on Aging, Grant Number T32-AG000186. A preliminary draft of this paper was circulated on April 16, 2018. Replication code is available at <http://economics.mit.edu/grad/lsun20/>. The companion Stata package `eventstudyweights` is available from the SSC repository.

[†]Department of Economics, MIT, 77 Massachusetts Avenue, Cambridge, MA 02139. Corresponding author; lsun20@mit.edu.

[‡]Cornerstone Research, 699 Boylston St, Boston, MA 02116. The views expressed herein are solely those of the author, who is responsible for the content, and do not necessarily represent the views of Cornerstone Research.

Assumption 3 requires that each cohort experiences the same path of treatment effects. Treatment effects need to be the same across cohorts in every relative period for homogeneity to hold, whereas for heterogeneity to occur, treatment effects only need to differ across cohorts in one relative period. The assumption of treatment effect homogeneity is therefore strong, and in Section 3.4.1, we describe how it can be violated in applied settings.

Our notion of treatment effect homogeneity does not preclude dynamic treatment effects; it only imposes that cohorts share the same path of treatment effects. The related literature sometimes formulates restrictions on the dynamics of treatment effects as another notion of treatment effect homogeneity. Athey and Imbens (2018) propose an assumption that “restricts the heterogeneity of the treatment effects over time,” which implies $CATT_{e,\ell}$ can vary over e but not over ℓ . Borusyak and Jaravel (2017) refer to one type of treatment effects heterogeneity as “only across the time horizon,” which implies $CATT_{e,\ell}$ can vary over ℓ but not over e . Callaway and Sant’Anna (2020a) allow for “arbitrary treatment effect heterogeneity” when $CATT_{e,\ell}$ varies across cohorts and over time. Similarly, de Chaisemartin and D’Haultfœuille (2020) describe treatment effects that may be “heterogeneous across groups and over time periods.” Goodman-Bacon (2018) allow heterogeneous effects to either “vary across units but not over time” or “vary over time but not across units.” The literature has not converged on a single notion of treatment effects heterogeneity with time-varying treatment. Since researchers are interested in dynamic treatment effects when using a “dynamic” specification, we do not restrict the path of treatment effects but rather use “heterogeneity” to describe variation across cohorts only.

2.3 Relevance in the applied literature

To gauge the empirical relevance of our results, we survey the estimation methods used by the twelve papers collected by Roth (2019) from three leading economics journals that contain the phrase “event study” in their main text.¹ From this sample of applied papers, we learn what specifications empirical researchers are actually using when estimating two-way fixed effects regressions. Four papers in this sample consider the simple setting where units either receive their first treatment at the same time or never receive the treatment. The other eight papers in this sample consider the more complex setting where treated units receive their

¹We follow the selection criteria in Roth (2019): the original sample consists of 70 total papers, but is further constrained to these twelve papers with publicly available data and code. The data and code are used to determine exactly the specification estimated in these papers.

for all ℓ and all t . We can represent the relative period bin indicator as

$$\mathbf{1}\{t - E_i \in g\} = \sum_{\ell \in g} \mathbf{1}\{t - E_i = \ell\} = \sum_{\ell \in g} D_{i,t}^\ell. \quad (5)$$

Static specification. For a “static” specification $\mathcal{G}$ contains a single element equal to $g = [0, T]$. The indicator $\mathbf{1}\{t - E_i \in g\}$ is equivalent to an indicator for whether unit i has received its initial treatment by t : $\mathbf{1}\{E_i \leq t\}$. The “static” specification thus takes the following form

$$Y_{i,t} = \alpha_i + \lambda_t + \mu_g \sum_{\ell \geq 0} D_{i,t}^\ell + v_{i,t} \quad (6)$$

and the corresponding set of excluded relative periods is $g^{excl} = [-T, -1]$.

Dynamic specification. “Dynamic” specifications encompass any specifications of $\mathcal{G}$ where $\mathcal{G}$ contains more than one element, thus treatment effects are allowed to vary over time non-parametrically. In its most flexible form, a “fully dynamic” specification takes the following form

$$Y_{i,t} = \alpha_i + \lambda_t + \sum_{\ell=-K}^{-2} \mu_\ell D_{i,t}^\ell + \sum_{\ell=0}^L \mu_\ell D_{i,t}^\ell + v_{i,t} \quad (7)$$

and the corresponding set of excluded relative periods is $g^{excl} = \{-T, \dots, -K-1, -1, L+1, \dots, T\}$.

Excluding some relative periods from the “fully dynamic” specification is necessary to avoid multi-collinearity, either among the relative period indicators $D_{i,t}^\ell$, or with the unit and time fixed effects. For example, when there are no never-treated units i.e. $\infty \notin \text{supp}(E_i)$ but with a panel balanced in calendar time, we need to exclude at least two relative period indicators in $\mathcal{G}$. These collinearities are discussed by Borusyak and Jaravel (2017): one multi-collinearity comes from the relative period indicators summing to one for every unit $\sum_{\ell \in [-T, T]} D_{i,t}^\ell = 1$, and the other multi-collinearity comes from the linear relationship between two-way fixed effects and the relative period indicators, namely $t - E_i = \ell$.

Excluding relative periods close to the initial treatment is common in practice. Normalizing relative to the period prior to treatment is the most common - six out of the eight papers we survey do so, as reflected in the above specification where we drop $D_{i,t}^{-1}$. The remaining two papers exclude $D_{i,t}^0$.

Excluding distant relative periods is however less common (only one of the eight papers we survey does so). Instead researchers “bin” or “trim” distant relative periods. For “binning”, researchers bin distant

each $CATT_{e,\ell}$ using an interacted two-way fixed effects regression, then estimate the weight $Pr\{E_i = e \mid E_i \in [-\ell, T - \ell]\}$ using their sample analogs. In the final step, we average over the cohort-specific estimates associated with relative period ℓ . This method has a similar flavor as the method proposed by Gibbons et al. (2019). They first use an interacted model to estimate the treatment effect for each fixed effect group; the resulting group-specific estimates are averaged to provide the ATE. Their method improves fixed effects regressions in a cross-sectional setting, and our method builds on theirs by improving two-way fixed effects regressions in a panel setting. We therefore follow their terminology in calling our alternative estimator an “interaction-weighted” estimator.

4.1 Interaction-weighted estimator

We describe the estimation procedure in three steps (with more detailed definitions stated in Definition 4 of Online Appendix B).

Step 1. We estimate $CATT_{e,\ell}$ using a linear two-way fixed effects specification that interacts relative period indicators with cohort indicators, excluding indicators for cohorts from some set C :

$$Y_{i,t} = \alpha_i + \lambda_t + \sum_{e \notin C} \sum_{\ell \neq -1} \delta_{e,\ell} (\mathbf{1}\{E_i = e\} \cdot D_{i,t}^\ell) + \epsilon_{i,t}. \quad (26)$$

The exact specification depends on the cohort shares for a given application. If there is a never-treated cohort, i.e. $\infty \in \text{supp}\{E_i\}$, then we may set $C = \{\infty\}$ and estimate regression (26) on all observations. If there are no never-treated units, i.e. $\infty \notin \text{supp}\{E_i\}$, then we may set $C = \{\max\{E_i\}\}$, i.e. the latest-treated cohort and estimate regression (26) on observations from $t = 0, \dots, \max\{E_i\} - 1$. Lastly, if there is a cohort that is always treated, i.e. $0 \in \text{supp}\{E_i\}$, then we need to exclude this cohort from estimation.

The coefficient estimator $\widehat{\delta}_{e,\ell}$ from regression (26) is a DID estimator for $CATT_{e,\ell}$ with particular choices of pre-periods and control cohorts. As DID is likely a familiar estimator for applied researchers, we separate the more in-depth discussion of its definition, choices of pre-periods, and choices of control cohorts in Section 4.2.

Step 2. We estimate the weights $Pr\{E_i = e \mid E_i \in [-\ell, T - \ell]\}$ by sample shares of each cohort in the relevant period(s) $\ell \in g$.

Step 3. To form our IW estimator, we take a weighted average of estimates for $CATT_{e,\ell}$ from Step 1

with weight estimates from step 2. More formally, the IW estimator is

$$\widehat{v}_g = \frac{1}{|g|} \sum_{\ell \in g} \sum_e \widehat{\delta}_{e,\ell} \widehat{Pr}\{E_i = e \mid E_i \in [-\ell, T - \ell]\} \quad (27)$$

where $\widehat{\delta}_{e,\ell}$ is returned from step 1 and $\widehat{Pr}\{E_i = e \mid E_i \in [-\ell, T - \ell]\}$ is the estimated weight returned from step 2. We normalize the weights further by the size of g . If g is a singleton, then its size is $|g| = 1$.

Validity of the IW estimator. Under the parallel trends and no anticipation assumptions the coefficient estimator $\widehat{\delta}_{e,\ell}$ from regression (26) is a consistent estimator for $CATT_{e,\ell}$. The sample shares of each cohort are also consistent estimators for the population shares. Thus, the IW estimator is consistent for a weighted average of $CATT_{e,\ell}$ with weights equal to the share of each cohort in the relevant period(s).

With a few standard assumptions (which we present as Assumption 4 in Online Appendix B) on regression (26), we can show that each IW estimator is asymptotically normal and derive its asymptotic variance. The large sample approximation allows us to estimate the variance of IW estimators directly without relying on bootstrapping as in Callaway and Sant'Anna (2020a). However, we only construct pointwise confidence interval valid for a given IW estimator $\widehat{v}_g$. The bootstrap-based inference by Callaway and Sant'Anna (2020a) constructs simultaneous confidence intervals that are valid for the entire path of $\widehat{v}_g$.

4.2 Difference-in-differences estimator for $CATT_{e,\ell}$

Definition 2. Assume cohort e is non-empty i.e. $\sum_{i=1}^N \mathbf{1}\{E_i = e\} > 0$. Assume there exists some pre-period $s < e$ and some set of control cohorts $C \subseteq \{c : e + \ell < c \leq T\}$ that are non-empty i.e. $\sum_{i=1}^N \mathbf{1}\{E_i \in C\} > 0$. Using the notion $\mathbb{E}_N$ to abbreviate the symbol $\frac{1}{N} \sum_{i=1}^N$, the DID estimator with pre-period s and control cohorts C estimates $CATT_{e,\ell}$ as

$$\widehat{\delta}_{e,\ell} = \frac{\mathbb{E}_N[(Y_{i,e+\ell} - Y_{i,s}) \cdot \mathbf{1}\{E_i = e\}]}{\mathbb{E}_N[\mathbf{1}\{E_i = e\}]} - \frac{\mathbb{E}_N[(Y_{i,e+\ell} - Y_{i,s}) \cdot \mathbf{1}\{E_i \in C\}]}{\mathbb{E}_N[\mathbf{1}\{E_i \in C\}]} \quad (28)$$

The assumption of non-empty cohort e , existence of pre-period and non-empty control cohorts makes the DID estimator well-defined. For example, DID estimators for cohort 0 are not well-defined because a pre-period does not exist for this cohort, which is why we exclude them from estimating regression (26) in Step 1.

Using the above definition, in regression (26) from Step 1 of our proposed method, the coefficient

estimator $\widehat{\delta}_{e,\ell}$ is a DID estimator for $CATT_{e,\ell}$ with pre-period $s = e - 1$ (because we exclude relative period $\ell = -1$) and some choice of control cohorts C . If there is a never-treated cohort i.e. $\infty \in \text{supp}\{E_i\}$, then we set the control cohort to be never-treated units $C = \{\infty\}$. If there are no never-treated units, i.e. $\infty \notin \text{supp}\{E_i\}$, then we set the control cohort $C = \{\max\{E_i\}\}$, i.e. the latest-treated cohort. Among all possible pre-periods, for regression (26) from step 1 of our proposed method we choose $s = e - 1$ and $C = \{\infty\}$ with never-treated units (or $\{\max\{E_i\}\}$ without never-treated units) because the resulting specification is a natural extension of the common specifications of two-way fixed effects regression.

Note that without never-treated units we need to drop time periods $t \geq \max\{E_i\}$ from estimating regression (26) because every unit will be treated in these periods. The DID estimators for $CATT_{e,\ell}$ for $e + \ell \geq \max\{E_i\}$ are thus not well-defined as the control cohort is empty $C = \emptyset$. For example, when there are just two cohorts with $E_i \in \{1, T\}$, one treated at $t = 1$ and the other treated in the last period, we then need to omit interaction terms involving the latest-treated cohorts as well as dropping observations from the last period $t = T$ from estimation.

Under some assumptions, this DID estimator $\widehat{\delta}_{e,\ell}$ is an unbiased and consistent estimator for $CATT_{e,\ell}$, a fact that we build on in deriving the probability limit of the IW estimator. We state this in the following proposition.

Proposition 5. *If Assumptions 1 and 2 hold, then the DID estimator using any pre-period $s < e$ and non-empty control cohorts C is an unbiased and consistent estimator for $CATT_{e,\ell}$.*

It is possible to relax the parallel trends assumption to allow the timing of treatment to depend on covariates. One can estimate $CATT_{e,\ell}$ consistently based on the inverse propensity score reweighted estimator proposed by Abadie (2005) and Callaway and Sant’Anna (2020a), the outcome regression approach by Heckman et al. (1997), and the doubly robust estimator recently proposed by Sant’Anna and Zhao (2020). The resulting estimates $\widehat{\delta}_{e,\ell}$ can then be plugged into step 3 to form our IW estimator. In particular, without covariates and for the case with never treated units, our approach coincides with Callaway and Sant’Anna (2020a). Therefore one can use the did R package developed by Callaway and Sant’Anna (2020b) to form the IW estimator.

The choice of the pre-period s and the control cohorts C depends on the trade-off between relaxing Assumption 1 or Assumption 2. If Assumption 2 is likely to hold, we can choose any pre-period $s < e$ and include not-yet-treated units in control cohorts $C = \{c : c > e + \ell\}$ as in Callaway and Sant’Anna (2020a). This