

ment positions, including job titles, start and end dates, and the employing firm.¹⁶

A key feature of the dataset is the standardized *seniority level* variable for each position, constructed by Revelio through an ensemble modeling approach based on multiple sources of information. This measure combines information from (i) the worker’s current job (title, firm, and industry), (ii) their work history (tenure and previous seniority), and (iii) their age. These three inputs produce separate scores, which are averaged and then categorized into seven standardized seniority levels: Entry Level, Junior Level, Associate Level, Manager Level, Director Level, Executive Level, and Senior Executive Level. In the analysis that follows, we group positions into two broad categories: *juniors* (Entry and Junior) and *seniors* (Associate and above).¹⁷ Supplemental Appendices A.1 and A.2 provide comprehensive data validation; Section 4.3 below summarizes their main findings.

We merge the worker résumé data with Revelio’s job postings database at the firm level. The job postings database tracks firms’ recruitment activity since September 2021 and contains, for each posting, a firm identifier, posting date, and the raw text of the job description. We use these job-posting texts to construct our measure of firm-level GenAI adoption, as described in Section 4.2.

In addition, we incorporate the occupational GenAI exposure measures from Eloundou et al. (2024), merged with our position-level data via O*NET SOC codes. We rely on the GPT-4–based beta exposure measure (as in, for example, Brynjolfsson et al., 2025a), and classify all the positions into three categories: low exposure (0–25th percentile), medium exposure (25–75th percentile), and high exposure (75–100th percentile).

We also construct a representative sample of *tasks* from Revelio’s job postings database. The extraction pipeline is motivated by Hampole et al. (2025), who construct firm-level AI exposure by mapping résumé text to O*NET tasks. From the Revelio job-postings corpus we draw a representative sample of 1,000 firms containing 355,013 postings, and map each posting’s free-text description to a subset of the 19,265 standardized O*NET tasks using a two-stage pipeline. The first stage is a TF-IDF retriever that returns the top 50 candidate O*NET tasks for each posting; the second stage is a large language model

¹⁶The initial data download was performed in May 2025. In April 2026, we re-downloaded the positions data to capture profile updates that users added with a lag, ensuring more complete coverage of recent employment spells.

¹⁷More details on Revelio’s seniority classification methodology are available at <https://www.data-dictionary.reveliolabs.com/methodology.html#seniority>.

tions and zero for low-exposure occupations; γ_{it} are firm-by-time fixed effects; and $\xi_{p(i)ct}$ are industry-by-exposure-by-time fixed effects. Equation 12 has the same structure as the seniority-based triple-difference in Equation 11, but here the third dimension is *occupational exposure* (high vs. low) rather than seniority. The coefficients β_j^s trace, separately for juniors and seniors, how employment in high-exposure occupations evolves relative to low-exposure occupations within the same firm and quarter, comparing adopters to non-adopters.

Figure 6 plots the estimated β_j^s coefficients for juniors and seniors, with both series normalized to zero in 2022Q4, the last pre-ChatGPT quarter. Two patterns stand out. First, throughout the pre-ChatGPT period, high-exposure occupations expand faster than low-exposure occupations at adopters relative to non-adopters for both seniority groups, consistent with adopters tilting their workforce toward exposed occupations over time. The two profiles track each other closely through 2022, with the senior coefficients rising somewhat faster than the junior coefficients. Second, beginning in 2023Q1, the two series diverge sharply. The junior series declines steadily, reaching a roughly 7 log point decline in high-exposure versus low-exposure junior employment at adopters by 2025Q1, relative to the 2022Q4 baseline. In contrast, the senior series continues to rise after 2022Q4, with no visible break around the ChatGPT launch.

Taken together, these results provide two insights. First, the post-2022 contraction in junior employment at adopters is concentrated in high-exposure occupations rather than reflecting a broad-based decline across all junior roles. Second, the design helps rule out a purely compositional explanation—namely, that high-exposure occupations declined economy-wide after 2023 and adopters simply employed a larger share of such roles. By comparing high- and low-exposure occupations *within* the same firm-by-quarter, and by running the same specification for seniors as a placebo-like comparison, we find that high-exposure occupations contract relative to low-exposure occupations only for juniors and only at adopters after the ChatGPT launch. The senior profile continues along its pre-existing upward trend, consistent with the interpretation that GenAI adoption is seniority-biased: it shifts demand away from junior workers in the very occupations where the technology is most applicable, while leaving senior employment in those same occupations essentially unaffected.

Supplemental Appendix A.3 reports a complementary heterogeneity exercise that re-

5.5 Decomposing the Decline in Junior Employment into Flows

The decline in junior employment in adopting firms can arise through three channels: (i) reduced hiring, (ii) increased separations, or (iii) increased promotions to senior positions. Our LinkedIn-based résumé data can be viewed as a detailed matched employer-employee dataset, allowing us to track these worker flows over time. Specifically, we estimate separate DiD regressions of the following form:

$$y_{it} = \alpha + \beta (\text{Adopt}_i \times \text{Post}_t) + \delta_t + \gamma_i + \xi_{pt} + \varepsilon_{it}, \quad (15)$$

where y_{it} denotes the number of junior hires, number of junior separations, or the junior promotion rate³¹ in firm i at quarter t . γ_i denotes firm fixed effects, δ_t are time fixed effects and ξ_{pt} are industry-by-time fixed effects.

Table 4 reports the estimated β from Equation 15. The results indicate that the sharp contraction in junior employment among adopters is driven primarily by a slowdown in hiring. Specifically, the coefficient on *Hiring* implies that, relative to non-adopters, GenAI-adopting firms hired on average 4.0 fewer junior workers per quarter after 2023Q1—about an 80 percent reduction relative to the pre-period mean of 5.06 hires per firm-quarter. Junior separations also declined among adopters, by roughly 1.1 workers per quarter, partially offsetting the hiring contraction; but this reduction is about one-fourth the magnitude of the hiring decline, so net junior employment still contracts sharply.³² Junior promotion rates rose by a small but statistically significant 0.033 percentage points (relative to a pre-period mean of 0.82 percent). Because promotions move workers out of the junior pool, this rise further reduces the junior stock, though its magnitude is very small relative to the hiring channel. Appendix A.28 explores sectoral heterogeneity in the decline in junior hiring, showing that the contraction is broad-based across industries and not driven disproportionately by any single sector.

³¹The junior promotion rate is defined as $\frac{\text{Promotions}_t}{\text{Junior Employment}_{t-1}}$, expressed in percentage points.

³²See Supplemental Appendix A.26 for the time-series DiD results for junior hires and exits.