

5 Validation of Measures

5.1 Comparison to Earlier Efforts

This paper aims to build on a number of previous empirical studies examining the occupational exposure to advances in AI and/or automation. Previous studies have used a variety of methods, including:

- Using occupational taxonomies like O*NET to characterize which occupations have routine vs. non-routine and manual vs. cognitive task content (Autor et al., 2003; Acemoglu and Autor, 2011a).
- Mapping text descriptions of tasks to descriptions of technological advances in patents. (Kogan et al., 2021; Webb, 2020)
- Linking capabilities of AI systems to occupational abilities and aggregating exposure estimates to the occupations where those abilities are required. (Felten et al., 2018, 2023)
- Mapping the results of AI task benchmark evaluations (ImageNet, Robocup, etc.) to 59 worker tasks through a set of 14 cognitive abilities drawn from the cognitive science literature. (Tolan et al., 2021)
- Expert labeling of automation potential for a set of O*NET occupations where experts had high confidence, combined with a probabilistic classifier to estimate automation potential for the remainder of O*NET occupations. (Frey and Osborne, 2017)
- Developing a rubric for evaluating the "suitability for machine learning" (SML) of activities that workers are completing in the economy (Brynjolfsson and Mitchell, 2017; Brynjolfsson et al., 2018, 2023).

We provide a set of summary statistics on many of these prior efforts in Table 8.

This paper's methodology primarily builds upon the SML approach by developing a rubric to evaluate the overlap between LLM capabilities and worker tasks as reported in the O*NET database. Table 9 presents the results of OLS regressions of our new LLM exposure measurements on occupation-level exposure measures from (Felten et al., 2018) ("AI Occupational Exposure Score" in the table), (Frey and Osborne, 2017) (Frey & Osborne Automation), scores from all three technologies in (Webb, 2020), normalized routine manual and cognitive scores from (Acemoglu and Autor, 2011a), and (Brynjolfsson et al., 2018, 2023) (SML). We also use annualized occupational salaries from the most recent BLS Occupational Employment Survey as a control. There are four separate output variables representing new scores in this paper that are predicted by earlier efforts.

GPT-4 Exposure Rating 1 corresponds to our overall exposure rubric as evaluated by GPT-4, where full exposure potential is coded as 1, no exposure potential is coded as 0, and partial exposure (E2 in our labeling scheme) is coded as 0.5. GPT-4 Exposure Rating 2 is scored similarly for overall exposure, but with a slightly different prompt. The results are very similar across the two prompts. Human Exposure Rating represents the same rubric as in GPT-4 Exposure Rating 1 but is scored by humans, as discussed in an earlier section of the paper. These results correspond to the β set of statistics presented above, with supplemental tasks having half the weight of core tasks within an occupation. These weights sum to one (core/supplemental distinctions are determined by O*NET).

The results across each type of measurement are consistent. We find generally positive and statistically significant correlations between our LLM exposure measures and previous measurements targeting software and AI. Encouragingly, the SML exposure scores by occupation show significant and positive associations with the exposure scores we develop in this paper, demonstrating a level of cohesion between the two studies with similar approaches. The Webb software and AI patent-based measures, SML, and normalized (demeaned