

Claude is mostly used for work, but use cases diversify with adoption

Our data, relying on a [privacy-preserving](#)¹ analysis of 1 million Claude.ai conversations,² reveals striking geographic differences in how Claude is adopted. Claude is predominantly used for work, across the globe and across the United States. However, there is geographic variation in use cases. At the global level, the Balkans and Brazil have the highest relative share of work use (see Figure 3.1), and Indonesia stands out with the highest share of coursework. At the US state level, New York stands out as the state using Claude relatively the most for work.




Figure 3.1: Share of work use of Claude.ai globally.

The share of conversations for a given country that are classified as work, as opposed to personal or coursework. The different tiers reflect a country's position within the global distribution of the Anthropic AI Usage Index as defined in chapter 1.^{3,4,5} We only include countries with at least 200 observations in our sample for this figure because of the uncertainty of the measure for low-usage countries in our random sample. The underlying data includes Claude.ai Free, Pro and Max usage.

Use case differences are related to a country's per capita income, which, in turn, is related to per capita AI adoption. We observe that work use cases and personal use cases of Claude are more common in higher income countries, while coursework use cases are more common in lower income countries (see Figure 3.2). Interestingly, these findings converge with [recent work by Microsoft](#) showing that AI use for school is associated with lower per capita income, whereas AI use for leisure is associated with higher per capita income.

preparing interpretive reports—have high success rates. These occupations have low task coverage because AI can’t do the hands-on or administrative work in their job profiles, but it succeeds on the core knowledge work that dominates their workday.

Microbiologists fall below the 45-degree line, suggesting lower effective AI coverage than would be predicted by task coverage alone. Claude covers half of their tasks, but not their most time-intensive: hands-on research using specialized lab equipment.

This measure arguably gives a more realistic picture of job-level AI penetration. However, its implications depend on how often these Claude conversations actually displace or augment work that would otherwise be done by humans. For data entry clerks, AI likely does substitute for tasks previously performed manually. But when a Claude conversation maps to a teacher performing a lecture, it is less clear how this translates to reduced lecture time on the job. In future work, we could leverage our 1P API data to understand which of these tasks are being integrated into production workflows.

AI’s impact on the task content of jobs

Beyond how much of a worker’s day AI can successfully perform, a separate question is which tasks get covered, and whether those tend to be the high-skill or low-skill components of the job. Recent research has studied changes in the task mix within jobs to understand AI’s impact on wages and employment ([Autor and Thompson 2025](#); [Hampole et al 2025](#)). A key insight is that automation’s effects depend not just on how many tasks are covered, but on which tasks.

To see how jobs change when we remove the tasks AI can perform, we first construct a measure of the level of skill required for each task. O*NET doesn’t provide task-level education requirements, so we train a model that predicts years of schooling from task embeddings, using the BLS’s occupation-level education as the target.⁴ This way, a low-education occupation may still have a high-skill task if it looks like those that tend to exist in high-education occupations. For example, Legal Secretaries is a 12-year education occupation, but the task “Review legal publications and perform database searches to