

Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects*

Liyang Sun[†] and Sarah Abraham[‡]

September 22, 2020

Abstract

To estimate the dynamic effects of an absorbing treatment, researchers often use two-way fixed effects regressions that include leads and lags of the treatment. We show that in settings with variation in treatment timing across units, the coefficient on a given lead or lag can be contaminated by effects from other periods, and apparent pretrends can arise solely from treatment effects heterogeneity. We propose an alternative estimator that is free of contamination, and illustrate the relative shortcomings of two-way fixed effects regressions with leads and lags through an empirical application.

Keywords: difference-in-differences, two-way fixed effects, pretrend test

*We are grateful to Isaiah Andrews, Amy Finkelstein, Anna Mikusheva, and Heidi Williams for their guidance and support. We thank Alberto Abadie, Jonathan Cohen, Nathan Hendren, Peter Hull, Guido Imbens, Yunan Ji, Sylvia Klosin, Kevin Kainan Li, Paichen Li, Therese A. McCarty, Whitney Newey, James Poterba, Pedro H. C. Sant’Anna, Gergely Ujhe-lyi and Helen Willis for helpful discussions. This research was supported by the National Institute on Aging, Grant Number T32-AG000186. A preliminary draft of this paper was circulated on April 16, 2018. Replication code is available at <http://economics.mit.edu/grad/lsun20/>. The companion Stata package `eventstudyweights` is available from the SSC repository.

[†]Department of Economics, MIT, 77 Massachusetts Avenue, Cambridge, MA 02139. Corresponding author; lsun20@mit.edu.

[‡]Cornerstone Research, 699 Boylston St, Boston, MA 02116. The views expressed herein are solely those of the author, who is responsible for the content, and do not necessarily represent the views of Cornerstone Research.

following differences. First, in the setting where there is no never-treated group, our method uses the last cohort to be treated as a control group, whereas Callaway and Sant’Anna (2020a) use the set of not-yet-treated cohorts. Our method and theirs thus rely on different, but non-nested parallel trends assumptions. Second, our estimation method can be cast as a regression specification and thus may be more familiar to applied researchers. However, a third difference is that the procedure of Callaway and Sant’Anna (2020a) allows for conditioning on time-varying covariates. de Chaisemartin and D’Haultfœuille (2020) and Goodman-Bacon (2018) respectively propose alternative estimators and diagnostic tools for estimation of causal effects in staggered settings, but do not consider the estimation of the dynamic path of treatment effects as we do.

2 Event studies design

In this section we first formalize the “event studies design”. As discussed in Section 2.3, based on how this term is deployed in the empirical literature, an *event study* design is a *staggered adoption* design where units are treated at different times, and there may or may not be never treated units. It also nests a *difference-in-differences* design, where units are either first treated at time t_0 or never treated.

Specifically, we consider a setting with a random sample of N units observed over $T + 1$ time periods, where T is fixed. For each $i \in \{0, \dots, N\}$ and $t \in \{0, \dots, T\}$, we observe the outcome $Y_{i,t}$ and treatment status $D_{i,t} \in \{0, 1\}$: $D_{i,t} = 1$ if i is treated in period t and $D_{i,t} = 0$ if i is not treated in period t . Throughout we assume that the observations $\{Y_{i,t}, D_{i,t}\}_{t=0}^T$ are independent and identically distributed (i.i.d.).

In the general case of event studies we focus on an *absorbing* treatment such that the treatment status over time is a non-decreasing sequence of zeros and then ones, i.e. $D_{i,s} \leq D_{i,t}$ for $s < t$. We can thus uniquely characterize a treatment path by the time period of the initial treatment, denoted with $E_i = \min \{t : D_{i,t} = 1\}$. If unit i is never treated i.e. $D_{i,t} = 0$ for all t , we set $E_i = \infty$. Based on when they first receive the treatment, we can also uniquely categorize units into disjoint cohorts e for $e \in \{0, \dots, T, \infty\}$, where units in cohort e are first treated at the same time $\{i : E_i = e\}$.

We define $Y_{i,t}^e$ to be the potential outcome in period t when unit i is first treated in time period e . We define $Y_{i,t}^\infty$ to be the potential outcome if unit i never receives the treatment, which we call the “baseline outcome”. Since the timing of the initial treatment uniquely characterizes one’s treatment path, we can