

firm regardless of AI exposure. For workers aged 22-25, we find a 15 log-point decline in relative employment for the most AI-exposed quintiles compared to the least exposed quintile, a large and statistically significant effect. Estimates for other age groups are much smaller in magnitude and not statistically significant. These findings imply that the employment trends we observe are not driven by differential shocks to firms that employ a disproportionate share of AI-exposed young workers.

Fifth, the labor market adjustments are visible in employment more than in compensation. In contrast to our findings for employment, we find little difference in annual salary trends by age or exposure quintile, suggesting possible wage stickiness. If so, AI may have larger effects on employment than on wages, at least initially, or even that AI may boost wages for as many workers as it hurts.

Sixth, the above facts are largely consistent across sample constructions designed to address various alternative explanations for the core findings. Our results are not driven solely by computer occupations or by occupations susceptible to remote work and outsourcing. We also find that the AI exposure taxonomy did not meaningfully predict employment outcomes for young workers further back in time, before the widespread use of LLMs, including during the unemployment spike driven by the COVID-19 pandemic. The patterns we observe in the data appear most acutely starting in late 2022 and early 2023, around the time of rapid proliferation of generative AI tools.³ They hold for both occupations with a high share of college graduates and ones with a low college share, suggesting deteriorating education outcomes during COVID-19 do not drive our results. For non-college workers, we find evidence that experience may serve as less of a buffer to labor market disruption, as low college share occupations exhibit divergent employment outcomes by AI exposure up to age 40.

While we explore a variety of alternative explanations, we caution that the facts we document may in part be influenced by factors other than generative AI. Taken as a whole, our results are consistent with the hypothesis that generative AI has begun to affect entry-level employment. We intend to continue to track the data on an ongoing basis to assess whether these trends change in the future.

³In particular, OpenAI introduced ChatGPT in November 2022. It was reported to have over 100 million active users by January 2023 and 1.7 billion website visitors in October 2023 (DeVon, 2023).

ADP services does not exactly match the distribution of firms across the broader US economy. Further details on differences in firm composition can be found in [Cajner et al. \(2018\)](#) and [ADP Research \(2025\)](#).¹³

3.2 Occupational AI Exposure

We use two approaches for measuring occupational exposure to AI. The first uses exposure measures from [Eloundou et al. \(2024\)](#), who estimate AI exposure by O*NET task using ChatGPT validated with human labeling. They construct occupational exposure measures by aggregating the task data to the 2018 SOC code level. We focus on the GPT-4 based β exposure measures from their paper.

The second approach we take uses generative AI usage data from the Anthropic Economic Index ([Handa et al., 2025](#)). This index reports the estimated share of queries pertaining to each O*NET task based on several million conversations with Claude, Anthropic’s generative AI model. It aggregates the data to the occupational level based on these task shares. One feature of the Anthropic Economic Index is that for each task it reports estimates of the share of queries pertaining to that task that are “automotive,” “augmentative,” or none of the above. We use this as an estimate of whether AI usage for an occupation is primarily complementary or substitutable with labor.¹⁴

We use a 2010 SOC code to 2018 SOC code crosswalk from the BLS to merge the exposure measures to the payroll data. Table [A1](#) shows example occupations for each AI exposure measure.

3.3 Other Data

To compare employment changes for teleworkable versus non-teleworkable occupations, we use data from [Dingel and Neiman \(2020\)](#). We use the Personal Consumption Expenditure index from the BLS to compute real earnings, indexed to October 2017. We use monthly Current Population Survey (CPS) data as a comparison for our main findings.

¹³[Cajner et al. \(2018\)](#) find a somewhat higher share of manufacturing and services firms compared to the QCEW using data from March 2016. They also find that ADP somewhat overrepresents firms in the Northeast. In addition, firms using ADP tend to grow faster on average than the typical firm in the US economy.

¹⁴[Handa et al. \(2025\)](#) use Claude to classify conversations into six categories. Directive and Feedback Loop conversations are considered Automotive; Task Iteration, Learning, and Validation are considered Augmentative. They also instruct the model to choose None of the above “liberally.” Table [A2](#) reproduces Table 1 from [Handa et al. \(2025\)](#) and shows more details about the automation and augmentation measures.

of whether the usage of AI in an occupation is primarily a substitute or complement for labor. Table A1 shows example occupations that are in the highest and lowest exposure category for each measure.

Panel A of Figure 3 shows employment changes by overall prevalence of related Claude queries for 22-25 year-olds.¹⁶ The patterns match the findings using the Eloundou et al. (2024) measures closely. Panel B likewise shows that the occupations with the highest estimated automation shares have experienced declining employment for the youngest workers.

In contrast, Panel C indicates that the occupations with the highest estimated augmentation shares have *not* experienced a similar pattern. Employment changes for young workers are not ordered by augmentation exposure, as the fifth quintile has among the fastest employment growth. The findings are consistent with automative uses of AI substituting for labor while augmentative uses do not.¹⁷

4.4 Fact 4: Employment declines for young, AI-exposed workers remain after conditioning on firm-time shocks

While our results so far are consistent with the hypothesis that generative AI is causing a decline in entry-level employment, there are plausible alternative explanations. One class of explanations is that our patterns are explained by industry- or firm-level shocks correlated with sorting patterns by age and measured AI exposure. For example, one possibility is that young workers with high measured AI exposure are disproportionately likely to sort to firms heavily susceptible to interest rate increases.

We test for a class of such confounders by controlling for a rich set of fixed effects. For each age group, we estimate the Poisson regression

$$\log(E[y_{f,q,t}]) = \sum_{q' \neq 1} \sum_{j \neq -1} \gamma_{q',j} 1\{t = j\} 1\{q' = q\} + \alpha_{f,q} + \beta_{f,t} + \epsilon_{f,q,t} \quad (4.1)$$

¹⁶See Figures A6 through A8 for other age groups.

¹⁷Occupations in the first two augmentation quintiles have very low Claude usage (0.01% and 0.09% of conversations for the average occupation, respectively), with a high share of conversations classified as neither automative nor augmentative. In contrast, occupations in the third through fifth quintiles average 0.47%, 0.39%, and 0.33% of Claude conversations. For the automation measure, overall Claude usage increases on average with the automation share, with the lowest exposure group averaging 0.05% of conversations and the highest group averaging 0.73%. Figures A9 and A10 show that automation and augmentation results are similar when dropping occupations with low overall Claude usage.

f indexes firms, q indexes Eloundou et al. (2024) exposure quintiles, and t indexes months, with $t = -1$ corresponding to October 2022. The outcome variable $y_{f,q,t}$ is employment in f, q, t . Equation 4.1 is a Poisson event study regression controlling for firm-quintile effects, $\alpha_{f,q}$, and firm-time effects, $\beta_{f,t}$. The firm-time effects absorb aggregate firm shocks that impact each exposure quintile equally. The firm-quintile effects adjust for baseline differences in hiring across quintiles within the firm. The coefficients of interest, $\gamma_{q,t}$, measure differential changes in employment growth across quintiles after accounting for firm-time effects and firm-quintile effects.¹⁸

We run this regression separately for each age group. For each regression, we restrict to firms that hire at least 10 workers within the age group in every period of the sample. Further, $\sum_t y_{f,q,t}$ must equal at least 100 for each q , meaning that the firm must at least employ on average about 2 workers from each exposure quintile across months in the sample.¹⁹ Standard errors are clustered by firm.

Results are in Figure 4, which plots the $\gamma_{q,t}$ coefficients for each age group. For workers aged 22-25, estimates for higher quintiles are large and statistically significant, with a 15 log point decline in relative employment comparable in magnitude to the estimates in the raw data in Figure 2. Estimates for other age groups are generally much smaller in magnitude and not statistically significant. These findings imply that the employment trends we observe are not driven by differential shocks to firms that employ a disproportionate share of AI-exposed young workers.

One alternative confounder that would not be controlled for with firm-time effects is that even conditional on the firm workers with high AI exposure were excessively hired after the COVID-19 pandemic, leading to a subsequent contraction in their hiring. To assess such alternatives we consider various other robustness checks in Section 4.6, such as removing computer occupations and conditioning on whether the occupation is amenable to work from home.

¹⁸Because of zero counts in the outcome variable, we estimate a Poisson regression instead of an OLS regression in logs following guidance from Chen and Roth (2024).

¹⁹Results are not sensitive to these restrictions, though there must be at least one non-zero value in each firm-month and each firm-quintile for observations to not get dropped in the Poisson regression.