


Each of the selected 608 tasks was then presented to a panel of experts. For each task, the experts were shown (i) the occupation and cluster information, (ii) the average survey score (rescaled to 0–1), and (iii) the full task description. If they agreed with the survey score, they left it unchanged; otherwise, they entered a revised expert score (0–1).

In many cases, the differences between the survey and expert scores are non-trivial. They also vary significantly across occupational groups, revealing distinct patterns based on the nature of pertaining tasks (Figure 13). In occupations dominated by manual and physical tasks, such as Craft and Related Trades Workers, Plant and Machine Operators, and Farmers, Gardeners, Foresters, and Fishers, experts consistently assigned lower scores compared to the survey respondents, reflecting the practical limitations of automating tasks that require physical manipulation or manual dexterity. In contrast, occupations involving more cognitive or knowledge-intensive tasks, such as Clerical Support Workers and Professionals, show a much more mixed picture. While some tasks in these groups received lower scores, many others were revised upwards by the expert group, reflecting the greater adaptability of generative AI for structured, data-driven, or communication-focused tasks. For example, Clerical Support Workers display a combination of increases and decreases, suggesting nuanced task-specific assessments where automation potential depends heavily on the nature of individual tasks.

► 5 Adjustment of survey sub-sample

In the next step, we compared the average of expert scores with the original mean survey scores using two independent AI models as impartial arbiters to propose an adjusted score.¹³ The models were informed about the sources of the scores in such a way that would prevent them from automatically siding with the expert group: group one was defined as *“practical knowledge of the work and types of tasks they were presented, since they worked in the same occupational groups from which we selected tasks they scored”*, while the second group was described as *“respondents with more theoretical knowledge of technology and labour markets”*. Rather than simply averaging the two scores, the models were explicitly instructed to focus their analysis on the task characteristics and the current capabilities of generative AI and required to provide justifications for each adjustment and its relative position to the original scores.

► Figure 13. Task-level scores from the survey and experts, compared to AI-arbitrated scores



The overall correlation between the adjusted scores produced by the two AI models is 0.96, with correlations within occupational groups exceeding 0.9 in all but one case¹⁴ (Figure 14). This high level of agreement demonstrates that, when provided with sufficient background knowledge in the form of reference ranges of scores, the models converge on remarkably similar predictions. Notably, only three tasks exhibited an absolute difference in scores exceeding 0.2 (orange dots in Figure 13). In these cases, GPT-4o tended to adopt a more optimistic perspective on the feasibility of automation, whereas Gemini emphasized the importance of human judgment as a critical factor.

¹³ We used OpenAI's GPT-4o and Google's Gemini Flash 1.5, which have comparable abilities. We also tested OpenAI o1, which at the time of writing was among the models with strongest capacities of « algorithmic reasoning » and observed very similar results. The process was set up as a cycle of sequential API calls, which loop over each task and capture the response into a data frame.

¹⁴ Public authorities, senior officials and managers in Figure 10.