

NBER WORKING PAPER SERIES

DIFFERENCE-IN-DIFFERENCES WITH VARIATION IN TREATMENT TIMING

Andrew Goodman-Bacon

Working Paper 25018

<http://www.nber.org/papers/w25018>

NATIONAL BUREAU OF ECONOMIC RESEARCH

1050 Massachusetts Avenue

Cambridge, MA 02138

September 2018

I thank Michael Anderson, Martha Bailey, Marianne Bitler, Brantly Callaway, Kitt Carpenter, Eric Chyn, Bill Collins, John DiNardo, Andrew Dustan, Federico Gutierrez, Brian Kovak, Emily Lawler, Doug Miller, Sayeh Nikpay, Pedro Sant'Anna, Jesse Shapiro, Gary Solon, Isaac Sorkin, Sarah West, and seminar participants at the Southern Economics Association, ASHEcon 2018, the University of California, Davis, University of Kentucky, University of Memphis, University of North Carolina Charlotte, the University of Pennsylvania, and Vanderbilt University. All errors are my own. The views expressed herein are those of the author and do not necessarily reflect the views of the National Bureau of Economic Research.

NBER working papers are circulated for discussion and comment purposes. They have not been peer-reviewed or been subject to the review by the NBER Board of Directors that accompanies official NBER publications.

© 2018 by Andrew Goodman-Bacon. All rights reserved. Short sections of text, not to exceed two paragraphs, may be quoted without explicit permission provided that full credit, including © notice, is given to the source.

Difference-in-differences (DD) is both the most common and the oldest quasi-experimental research design, dating back to Snow’s (1855) analysis of a London cholera outbreak.¹ A DD estimate is the difference between the change in outcomes before and after a treatment (difference one) in a treatment versus control group (difference two): $(\bar{y}_{TREAT}^{POST} - \bar{y}_{TREAT}^{PRE}) - (\bar{y}_{CONTROL}^{POST} - \bar{y}_{CONTROL}^{PRE})$. That simple quantity also equals the estimated coefficient on the interaction of a treatment group dummy and a post-treatment period dummy in the following regression:

$$y_{it} = \gamma + \gamma_i TREAT_i + \gamma_t POST_t + \beta^{2x2} TREAT_i \times POST_t + u_{it} \quad (1)$$

The elegance of DD makes it clear which comparisons generate the estimate, what leads to bias, and how to test the design. The expression in terms of sample means connects the regression to potential outcomes and shows that, under a common trends assumption, a two-group/two-period (2x2) DD identifies the average treatment effect on the treated. All econometrics textbooks and survey articles describe this structure,² and recent methodological extensions build on it.³

Most DD applications diverge from this 2x2 set up though because treatments usually occur at different times.⁴ The processes that generate treatment variables naturally lead to variation in timing. Local governments change policy. Jurisdictions hand down legal rulings. Natural disasters strike across seasons. Firms lay off workers. In this case researchers estimate a regression with dummies for cross-sectional units (α_i) and time periods (α_t), and a treatment dummy (D_{it}):

$$y_{it} = \alpha_i + \alpha_t + \beta^{DD} D_{it} + e_{it} \quad (2)$$

¹ A search from 2012 forward of nber.org, for example, yields 430 results for “difference-in-differences”, 360 for “randomization” AND “experiment” AND “trial”, and 277 for “regression discontinuity” OR “regression kink”.

² This includes, but is not limited to, Angrist and Krueger (1999), Angrist and Pischke (2009), Heckman, Lalonde, and Smith (1999), Meyer (1995), Cameron and Trivedi (2005), Wooldridge (2010).

³ Inverse propensity score reweighting: Abadie (2005), synthetic control: Abadie, Diamond, and Hainmueller (2010), changes-in-changes: Athey and Imbens (2006), quantile treatment effects: Callaway, Li, and Oka (forthcoming).

⁴ Half of the 93 DD papers published in 2014/2015 in 5 general interest or field journals had variation in timing.

In contrast to our substantial understanding of the canonical 2x2 DD model, we know relatively little about the two-way fixed effects DD model when treatment timing varies. We do not know precisely how it compares mean outcomes across groups.⁵ We typically rely on general descriptions of the identifying assumption like “interventions must be as good as random, conditional on time and group fixed effects” (Bertrand, Duflo, and Mullainathan 2004, p. 250), and consequently lack well-defined strategies to test the validity of the DD design with timing. We have limited understanding of the treatment effect parameter that regression DD identifies. Finally, we often cannot evaluate when alternative specifications will work or why they change estimates.⁶

This paper shows that the two-way fixed effects DD estimator in (2) is a weighted average of all possible 2x2 DD estimators that compare timing groups to each other. Some use units treated at a particular time as the treatment group and untreated units as the control group. Some compare units treated at two different times, using the later group as a control before its treatment begins and then the earlier group as a control after its treatment begins. As in any least squares estimator, the weights on the 2x2 DD’s are proportional to group sizes *and* the variance of the treatment dummy within each pair. Treatment variance is highest for groups treated in the middle of the panel and lowest for groups treated at the extremes. This result clarifies the theoretical interpretation and identifying assumptions of the general DD model and creates simple new tools for describing the design and analyzing problems that arise in practice.

By decomposing the DD estimator into its sources of variation (the 2x2 DD’s) and providing an explicit interpretation of the weights in terms of treatment variances, my results

⁵ Imai, Kim, and Wang (2018) note “It is well known that the standard DiD estimator is numerically equivalent to the linear two-way fixed effects regression estimator if there are two time periods and the treatment is administered to some units only in the second time period. Unfortunately, this equivalence result does not generalize to the multi-period DiD design...Nevertheless, researchers often motivate the use of the two-way fixed effects estimator by referring to the DiD design (e.g., Angrist and Pischke, 2009).”

⁶ This often leads to sharp disagreements. See Neumark, Salas, and Wascher (2014) on unit-specific linear trends, Lee and Solon (2011) on weighting and outcome transformations, and Shore-Sheppard (2009) on age time fixed effects.

Equation (17) also shows exactly how to weight together averages of x_{it} and perform a single t -test that directly captures the identifying assumption. Generate a dummy for the effective treatment group, $B_k = w_k^T - w_k^C > 0$, then regress timing-group means, $\bar{x}_k$, on B_k weighting by $|w_k^T - w_k^C|$. The coefficient on B_k equals covariate differences weighted by the actual identifying variation, and its t -statistic tests the null of reweighted balance in (17). One can also use this strategy to test for pre-treatment trends in confounders (or the outcome) by regressing $\bar{x}_{kt}$ on B_k , year dummies, and their interaction, or the interaction of B_k with a linear trend using dates before any treatment starts.²⁶

The reweighted balance test has advantages over other strategies for testing balance in this setting. Regressing x_{it} on a constant and dummies for timing groups allows a test of the null of joint balance across groups ($H_0: \bar{x}_k - \bar{x}_U = 0, \forall k \in K$). With many timing groups, however, this F -test will have low power, it does not reflect the importance of each group ($w_k^T - w_k^C$), and does not show the sign or magnitude of any imbalance.²⁷ The reweighted test, on the other hand, has higher power and describes the sign and magnitude of imbalance. It is also better than a test for linear relationships between x_{it} and t_k^* or tests for balance between “early” and “late” treated units (Almond, Hoynes, and Schanzenbach 2011, Bailey and Goodman-Bacon 2015). Because the effective control group can include *both* the earliest and latest treated units, failing to find a linear relationship between $\bar{x}_k$ and t_k^* can miss the relevant imbalance between the most important “middle” units and the end points.

²⁶ Plotting confounders or pre-trends across groups, however, is important to ensure that a failure to reject does not reflect offsetting trends or covariate means across timing groups.

²⁷ Note that using an event-study specification to evaluate pre-trends does not test for joint equality of pre-trends either. It tests for common trends using the (heretofore unarticulated) estimator itself.