


tal based on comparative advantage. When automation enables capital to perform tasks previously done by labor more cheaply, it creates a *displacement effect* that reduces labor demand. Conversely, the creation of new tasks where labor maintains comparative advantage generates a *reinstatement effect* that increases labor demand (Acemoglu and Restrepo, 2019; Restrepo, 2024). For LLMs, which can automate cognitive tasks previously considered non-routine, the framework predicts that displacement will dominate in the short run, particularly for entry-level positions involving routine information processing and basic analysis. Senior roles requiring tacit knowledge, judgment, and complex interpersonal coordination should prove less vulnerable. This is consistent with our findings.

Our findings contribute to the emerging international evidence on labor market impacts of AI (del Rio-Chanona et al., 2025). In the United States, Brynjolfsson, Chandar and Chen (2025); Lichtinger and Hosseini Maasoum (2025); Dominski and Lee (2025) document employment declines in exposed industries concentrated in early-career positions. However, Chandar (2025) finds no employment effects using CPS data, and Humlum and Vestergaard (2025) report null results from Danish administrative data. These disparities may reflect several factors: differences in labor market institutions between countries, variation in the timing of LLM adoption across regions, and methodological choices regarding exposure measurement and outcome variables. Our UK evidence suggests that effects manifest differently across adjustment margins, with hiring responding more quickly and strongly than employment stocks.

Methodologically, our paper builds on recent work that also uses Revelio Labs data to construct firm-level AI exposure measures. Eisfeldt et al. (2023) show that higher-exposure US firms experience excess stock returns following the release of ChatGPT, while Hampole et al. (2025) develop task-level exposure measures varying across firms and occupations to study US labor demand between 2010 and 2023. We extend their work in three directions: we shift focus to the UK labor market, examine realized employment and job posting outcomes (as opposed to Eisfeldt et al. (2023)’s focus on financial data), and exploit data through mid-2025 to capture labor market adjustment to generative AI adoption (as opposed to Hampole et al. (2025)’s focus on earlier waves of AI).

In the UK context, we extend Henseke et al. (2025), who develop a Generative AI Susceptibility Index and provide suggestive evidence that wage premia decline

Figure 1: Effect of LLM exposure on firm-level jobs



Notes: This figure plots event study estimates of the effect of LLM exposure on the total number of employees in UK firms (log). Coefficients are estimated from Equation (3). All regressions include firm fixed effects and industry-by-period fixed effects. Standard errors are clustered at the firm level. Vertical line indicates November 2022, when ChatGPT was publicly released. Shaded areas denote 95% confidence intervals.

show up in aggregate employment statistics. High and low LLM exposure exhibit similar employment trajectories prior to November 2022, which provides support for the validity of the parallel trends assumption. When averaged over the entire post-treatment period (Table 3), a one standard deviation increase in LLM exposure is associated with a 0.3 percent reduction in total employment, although the aggregate effect is only marginally significant due to the delayed onset of impacts.

We next examine how these aggregate effects vary across different types of workers. Figure 2 decomposes the effects between high- and low-seniority employees. Panel A shows junior positions decline immediately and statistically significantly, while Panel B demonstrates that senior positions initially remain stable before declining gradually, with the decrease only becoming significant after 20 months. Panel C shows that the timing difference in the emergence of effects leads to a temporary shift toward senior employees, although the compositional change attenuates over time, with the seniority distribution gradually reverting toward pre-treatment levels by the end of our observation period. All three panels show parallel pre-trends.

Table 4: Effect of LLM exposure on occupation-level job postings

	Job postings (log)	Salary	Salary (log)	Salary (SD)
LLM exposure x Post	-0.039** (0.018)	-495.483*** (135.270)	-0.011*** (0.003)	-56.802 (108.358)
Occupation fixed-effects	Yes	Yes	Yes	Yes
Year-by-month fixed effects	Yes	Yes	Yes	Yes
Observations	16,815	16,815	16,815	16,804
Adjusted R ²	0.962	0.948	0.953	0.729

Notes: This table reports difference-in-differences estimates of the effect of LLM exposure on occupation-level outcomes. *Job postings (log)* is the log number of new vacancies in occupation o in month t . *Salary* is the mean advertised annual compensation, *Salary (log)* is its logarithm, and *Salary (SD)* is the within-occupation standard deviation of posted salaries. *Exposure* is the occupation-level LLM exposure score from Eloundou et al. (2024). *Post* equals one for November 2022 onward. Standard errors are clustered at the occupation level. Significance is denoted by * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

4.2.2 Occupation-Level Postings

This analysis examines occupation-level job posting data to understand how aggregate labor demand evolves for specific roles. This analysis exploits variation in LLM exposure across 384 occupations, from minimally affected roles such as cooks (exposure score: 0) to highly exposed positions like proofreaders and copy markers (exposure score: 0.86). We examine three outcomes: posting volumes to measure aggregate demand, advertised salaries to assess compensation responses, and within-occupation salary dispersion to measure wage polarization.

Figure 6 presents event study estimates for job posting volumes. Prior to November 2022, high- and low-exposure occupations follow parallel trajectories, validating our identification strategy. Following ChatGPT’s release, however, exposed occupations experience a sustained reduction in hiring demand, with effects first becoming statistically significant after 7 months and intensifying thereafter. After 20 months, we see a slight resurgence, and job postings converge back to their original levels. When averaged across the post-treatment period (Table 4), a one standard deviation increase in LLM exposure is associated with a 3.9% reduction in posting volume ($p = 0.029$). This is an economically meaningful change.

Turning to compensation, we observe immediate and persistent wage pressure for exposed occupations. Figure 7 indicates that posted salaries begin declining shortly after ChatGPT’s introduction. A one standard deviation increase in LLM exposure is associated with a £495 decline in posted salaries and a 1.1% reduction in log