





(a) Junior positions (Seniority 1–2)



(b) Senior positions (Seniority 3+)

Figure A.7: Keyword Frequency by Seniority Group, Pre-2023 vs. 2023+

where $\text{Share}_{f,t}^{\text{match}}$ is the share of positions at firm f in quarter t whose raw title contains a matching keyword (junior-leaning for juniors, senior-leaning for seniors). The regression is weighted by the number of positions per firm-quarter cell, and standard errors are clustered at the firm level. Figure A.9 reports the estimated coefficients. For juniors, the




Figure A.8: Share of Positions with Matching Keywords, by Seniority and Adoption

Notes: This figure plots, for each quarter, the share of positions whose raw title contains a “matching” keyword: junior-leaning keywords among junior-classified positions (left panel) and senior-leaning keywords among senior-classified positions (right panel), at GenAI-adopting (solid) and non-adopting (dashed) firms. Shares are position-weighted within each firm-quarter cell, then aggregated across firms by adoption status. The vertical dashed line marks the launch of ChatGPT in November 2022. Sample: quarters 2020Q1–2025Q1.

coefficients hover around zero throughout the sample, with a visible quarterly saw-tooth that reflects the seasonal pattern of summer hiring and no discernible break at the ChatGPT launch. For seniors, the coefficients drift mildly upward in the post-period, reaching approximately +0.3 percentage points by 2025Q1—a quantitatively small movement relative to a baseline share of roughly 70%. Both patterns suggest that the seniority classification remains essentially stable across the diffusion of GenAI.




Figure A.9: Event Study: Differential Evolution of Matching-Keyword Shares

Notes: This figure plots quarterly event-study coefficients β_k from the specification $\text{Share}_{f,t}^{\text{match}} = \alpha_f + \alpha_t + \sum_{k \neq 2022Q4} \beta_k \cdot \mathbf{1}\{t = k\} \cdot \text{Adopter}_f + \varepsilon_{f,t}$, estimated separately for juniors (left panel: share of positions with any junior-leaning keyword in title) and seniors (right panel: share with any senior-leaning keyword). Regressions are weighted by the number of positions per firm-quarter cell, and standard errors are clustered at the firm level. Shaded areas denote 95% confidence intervals. The vertical dashed line marks the reference quarter 2022Q4 (the last pre-period quarter). Sample: quarters 2020Q1–2025Q1.

A.2.2 Worker Experience by Seniority and Adoption Status

A potential concern with our seniority-based analysis is that the classification of positions may itself shift with GenAI diffusion—for instance, if adopting firms differentially relabel roles as more senior after 2022. If so, the decline in junior employment at adopting firms could partly reflect mechanical reclassification rather than genuine changes in labor demand. To assess this possibility, we examine the evolution of actual worker experience among juniors and seniors, separately at adopting and non-adopting firms. We measure experience as *Labor Market Experience* (LME): the number of years since the worker’s first observed position. For each firm-quarter-seniority cell, we compute the worker-weighted average LME and then aggregate across firms by adoption status.

Figure A.10 plots average LME over time, separately for juniors (left panel) and seniors (right panel). For each seniority group, the adopter and non-adopter series evolve in close parallel. The pre-existing level gap between adopters and non-adopters—about

0.56 years for juniors and 0.28 years for seniors at 2022Q4—remains roughly constant throughout the sample, with no visible break around the ChatGPT launch. The upward trend in each series reflects the gradual aging of the observed workforce: average LME rises by approximately 2.8 years in all four series over the five-year window.




Figure A.10: Average Labor Market Experience by Seniority and GenAI Adoption Status

Notes: This figure plots average Labor Market Experience (LME), defined as years since the worker’s first observed position, for juniors (left panel) and seniors (right panel) at GenAI-adopting (solid) and non-adopting (dashed) firms over time. Averages are worker-weighted within each firm-quarter cell and then aggregated across firms by adoption status. The vertical dashed line marks the launch of ChatGPT in November 2022. Sample: quarters 2020Q1–2025Q1.

To isolate differential dynamics from the common time trend, we estimate, separately for juniors and seniors, the specification

$$\overline{\text{LME}}_{f,t} = \alpha_f + \alpha_t + \sum_{k \neq 2022\text{Q4}} \beta_k \cdot \mathbf{1}\{t = k\} \cdot \text{Adopter}_f + \varepsilon_{f,t},$$

where f indexes firms and t indexes quarters. The regression is weighted by the number of workers in each firm-quarter cell, and standard errors are clustered at the firm level. Each coefficient β_k captures the adopter–non-adopter difference in average LME at quarter k , relative to the baseline quarter 2022Q4.

Figure A.11 reports the results. For juniors, the coefficients are tightly centered around zero throughout the sample, with no discernible break at the ChatGPT launch. For se-